

```
#importing the dependencies
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn.datasets
from sklearn.model_selection import train_test_split
from xgboost import XGBRegressor
from sklearn import metrics
```

```
pip install scikit-learn==1.1.3
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: scikit-learn==1.1.3 in /usr/local/lib/python3.9/dist-packages (1.1.3)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.9/dist-packages (from scikit-learn==1.1.3)
Requirement already satisfied: numpy>=1.17.3 in /usr/local/lib/python3.9/dist-packages (from scikit-learn==1.1.3)
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.9/dist-packages (from scikit-learn==1.1.3)
Requirement already satisfied: joblib>=1.0.0 in /usr/local/lib/python3.9/dist-packages (from scikit-learn==1.1.3)
```

Double-click (or enter) to edit

Double-click (or enter) to edit

```
#Importing housing data set
house_price_dataset = sklearn.datasets.load_boston()
```

```
/usr/local/lib/python3.9/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function load_boston
```

The Boston housing prices dataset has an ethical problem. You can refer to the documentation of this function for further details.

The scikit-learn maintainers therefore strongly discourage the use of this dataset unless the purpose of the code is to study and educate about ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original source::

```
import pandas as pd
import numpy as np

data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
data = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]
```

Alternative datasets include the California housing dataset (i.e. `sklearn.datasets.fetch_california_housing`) and the Ames housing dataset. You can load the datasets as follows::

```
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()
```

for the California housing dataset and::

```
from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)
```

```
for the Ames housing dataset.
warnings.warn(msg, category=FutureWarning)
```

```
print(house_price_dataset)
```

```
[2.7310e-02, 0.0000e+00, 7.0700e+00, ..., 1.7800e+01, 3.9690e+02,
 9.1400e+00],
[2.7290e-02, 0.0000e+00, 7.0700e+00, ..., 1.7800e+01, 3.9283e+02,
 4.0300e+00],
...,
[6.0760e-02, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9690e+02,
 5.6400e+00],
[1.0959e-01, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9345e+02,
 6.4800e+00],
[4.7410e-02, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9690e+02,
 7.8800e+00]], 'target': array([24. , 21.6, 34.7, 33.4, 36.2, 28.7, 22.9, 27.1, 16.5, 18.9,
18.9, 21.7, 20.4, 18.2, 19.9, 23.1, 17.5, 20.2, 18.2, 13.6, 19.6,
15.2, 14.5, 15.6, 13.9, 16.6, 14.8, 18.4, 21. , 12.7, 14.5, 13.2,
13.1, 13.5, 18.9, 20. , 21. , 24.7, 30.8, 34.9, 26.6, 25.3, 24.7,
21.2, 19.3, 20. , 16.6, 14.4, 19.4, 19.7, 20.5, 25. , 23.4, 18.9,
35.4, 24.7, 31.6, 23.3, 19.6, 18.7, 16. , 22.2, 25. , 33. , 23.5,
19.4, 22. , 17.4, 20.9, 24.2, 21.7, 22.8, 23.4, 24.1, 21.4, 20. ,
20.8, 21.2, 20.3, 28. , 23.9, 24.8, 22.9, 23.9, 26.6, 22.5, 22.2,
23.6, 28.7, 22.6, 22. , 22.9, 25. , 20.6, 28.4, 21.4, 38.7, 43.8,
33.2, 27.5, 26.5, 18.6, 19.3, 20.1, 19.5, 19.5, 20.4, 19.8, 19.4,
21.7, 22.8, 18.8, 18.7, 18.5, 18.3, 21.2, 19.2, 20.4, 19.3, 22. ,
20.3, 20.5, 17.3, 18.8, 21.4, 15.7, 16.2, 18. , 14.3, 19.2, 19.6,
23. , 18.4, 15.6, 18.1, 17.4, 17.1, 13.3, 17.8, 14. , 14.4, 13.4,
15.6, 11.8, 13.8, 15.6, 14.6, 17.8, 15.4, 21.5, 19.6, 15.3, 19.4,
17. , 15.6, 13.1, 41.3, 24.3, 23.3, 27. , 50. , 50. , 50. , 22.7,
25. , 50. , 23.8, 23.8, 22.3, 17.4, 19.1, 23.1, 23.6, 22.6, 29.4,
23.2, 24.6, 29.9, 37.2, 39.8, 36.2, 37.9, 32.5, 26.4, 29.6, 50. ,
32. , 29.8, 34.9, 37. , 30.5, 36.4, 31.1, 29.1, 50. , 33.3, 30.3,
34.6, 34.9, 32.9, 24.1, 42.3, 48.5, 50. , 22.6, 24.4, 22.5, 24.4,
20. , 21.7, 19.3, 22.4, 28.1, 23.7, 25. , 23.3, 28.7, 21.5, 23. ,
26.7, 21.7, 27.5, 30.1, 44.8, 50. , 37.6, 31.6, 46.7, 31.5, 24.3,
31.7, 41.7, 48.3, 29. , 24. , 25.1, 31.5, 23.7, 23.3, 22. , 20.1,
22.2, 23.7, 17.6, 18.5, 24.3, 20.5, 24.5, 26.2, 24.4, 24.8, 29.6,
42.8, 21.9, 20.9, 44. , 50. , 36. , 30.1, 33.8, 43.1, 48.8, 31. ,
36.5, 22.8, 30.7, 50. , 43.5, 20.7, 21.1, 25.2, 24.4, 35.2, 32.4,
32. , 33.2, 33.1, 29.1, 35.1, 45.4, 35.4, 46. , 50. , 32.2, 22. ,
20.1, 23.2, 22.3, 24.8, 28.5, 37.3, 27.9, 23.9, 21.7, 28.6, 27.1,
20.3, 22.5, 29. , 24.8, 22. , 26.4, 33.1, 36.1, 28.4, 33.4, 28.2,
22.8, 20.3, 16.1, 22.1, 19.4, 21.6, 23.8, 16.2, 17.8, 19.8, 23.1,
21. , 23.8, 23.1, 20.4, 18.5, 25. , 24.6, 23. , 22.2, 19.3, 22.6,
19.8, 17.1, 19.4, 22.2, 20.7, 21.1, 19.5, 18.5, 20.6, 19. , 18.7,
32.7, 16.5, 23.9, 31.2, 17.5, 17.2, 23.1, 24.5, 26.6, 22.9, 24.1,
18.6, 30.1, 18.2, 20.6, 17.8, 21.7, 22.7, 22.6, 25. , 19.9, 20.8,
16.8, 21.9, 27.5, 21.9, 23.1, 50. , 50. , 50. , 50. , 50. , 13.8,
13.8, 15. , 13.9, 13.3, 13.1, 10.2, 10.4, 10.9, 11.3, 12.3, 8.8,
7.2, 10.5, 7.4, 10.2, 11.5, 15.1, 23.2, 9.7, 13.8, 12.7, 13.1,
12.5, 8.5, 5. , 6.3, 5.6, 7.2, 12.1, 8.3, 8.5, 5. , 11.9,
27.9, 17.2, 27.5, 15. , 17.2, 17.9, 16.3, 7. , 7.2, 7.5, 10.4,
8.8, 8.4, 16.7, 14.2, 20.8, 13.4, 11.7, 8.3, 10.2, 10.9, 11. ,
```

```
15.2, 16.1, 17.8, 14.9, 14.1, 12.7, 13.5, 14.9, 20. , 16.4, 17.7,
19.5, 20.2, 21.4, 19.9, 19. , 19.1, 19.1, 20.1, 19.9, 19.6, 23.2,
29.8, 13.8, 13.3, 16.7, 12. , 14.6, 21.4, 23. , 23.7, 25. , 21.8,
20.6, 21.2, 19.1, 20.6, 15.2, 7. , 8.1, 13.6, 20.1, 21.8, 24.5,
```

```
#Pandas Data frame
```

```
house_price_dataframe = pd.DataFrame(house_price_dataset.data ,columns=house_price_dataset.feature_names)
```

```
#Print First 5 rows
```

```
house_price_dataframe.head()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90

```
#add the target column to the DataFrame (Price of house)
```

```
house_price_dataframe['price'] =house_price_dataset.target
```

```
house_price_dataframe.head()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90

```
# Checking No of rows and columns in the data frame
```

```
house_price_dataframe.shape
```

```
(506, 14)
```

```
#Check for missing values
```

```
house_price_dataframe.isnull().sum()
```

```
CRIM      0
ZN        0
INDUS     0
CHAS      0
NOX       0
RM        0
AGE       0
DIS       0
```

```

RAD      0
TAX      0
PTRATIO  0
B        0
LSTAT    0
price    0
dtype: int64

```

```

#statistical measures of the dataset
house_price_dataframe.describe()

```

	CRIM	ZN	INDUS	CHAS	NOX	RM	
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.00
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.57
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.14
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.90
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.02
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.50
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.07
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.00



```

#correlation between data
correlation = house_price_dataframe.corr()

```

```

#construction a heatmap


```

```

plt.figure(figsize=(10,10))
sns.heatmap(correlation,cbar=True, square =True, fmt='.1f',annot=True, annot_kws={'size':8},cmap='Blues')

```

<Axes: >



Splitting the data and Target

x= house_price_dataframe.drop(['price'],axis=1)

y= house_price_dataframe['price']

print(x)

print(y)

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	
..	
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	

	PTRATIO	B	LSTAT
0	15.3	396.90	4.98
1	17.8	396.90	9.14
2	17.8	392.83	4.03
3	18.7	394.63	2.94
4	18.7	396.90	5.33
..
501	21.0	391.99	9.67
502	21.0	396.90	9.08
503	21.0	396.90	5.64
504	21.0	393.45	6.48
505	21.0	396.90	7.88

[506 rows x 13 columns]

```

0      24.0
1      21.6
2      34.7
3      33.4
4      36.2
...
501    22.4
502    20.6
503    23.9
504    22.0
505    11.9
Name: price, Length: 506, dtype: float64

```

```
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_state=2)
```

```
print(x.shape,x_train.shape,x_test.shape)
```

```
(506, 13) (404, 13) (102, 13)
```

```
#Training the model
```

```
#XGBoost
```

```
model=XGBRegressor()
```

```
#training the model with x_train
```

```
model.fit(x_train,y_train)
```

```

XGBRegressor
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              n_estimators=100, n_jobs=None, num_parallel_tree=None,
              predictor=None, random_state=None, ...)

```

```
#prediction on trainning data
```

```
#Accuracy for prediction on trainning data
```

```
training_data_prediction=model.predict(x_train)
```

```
print(training_data_prediction)
```

```

[23.147501  20.99463  20.090284  34.69053  13.903663  13.510157
 21.998634  15.1940975 10.899711  22.709627  13.832816  5.592794
 29.810236  49.99096  34.89215  20.607384  23.351097  19.23555
 32.695698  19.641418 26.991022  8.401829  46.00729  21.708961
 27.062933  19.321356 19.288303  24.809872  22.61626  31.70493
 18.542515  8.697379  17.395294  23.700663  13.304856  10.492197
 12.688369  25.016556 19.67495  14.902088  24.193798  25.007143
 14.900281  16.995798 15.6009035 12.699232  24.51537  14.999952
 50.00104  17.525454 21.184624  31.998049  15.613355  22.89754

```

19.325378	18.717896	23.301125	37.222923	30.09486	33.102703
21.00072	49.999332	13.405827	5.0280113	16.492886	8.405072
28.64328	19.499939	20.586452	45.402164	39.79833	33.407326
19.83506	33.406372	25.271482	50.001534	12.521657	17.457413
18.61758	22.602625	50.002117	23.801117	23.317268	23.087355
41.700035	16.119293	31.620516	36.069206	7.0022025	20.3827
19.996452	11.986318	25.023014	49.970123	37.881588	23.123034
41.292133	17.596548	16.305374	30.034231	22.860699	19.810343
17.098848	18.898268	18.96717	22.606049	23.141363	33.183487
15.010934	11.693824	18.78828	20.80524	17.99983	19.68991
50.00332	17.207317	16.404053	17.520426	14.593481	33.110855
14.508482	43.821655	34.939106	20.381636	14.655634	8.094332
11.7662115	11.846876	18.69599	6.314154	23.983706	13.084503
19.603905	49.989143	22.300608	18.930315	31.197134	20.69645
32.21111	36.15102	14.240763	15.698188	49.99381	20.423601
16.184978	13.409128	50.01321	31.602146	12.271495	19.219482
29.794909	31.536846	22.798779	10.189648	24.08648	23.710463
21.991894	13.802495	28.420696	33.181534	13.105958	18.988266
26.576572	36.967175	30.794083	22.77071	10.201246	22.213818
24.483162	36.178806	23.09194	20.097307	19.470194	10.786644
22.671095	19.502405	20.109184	9.611871	42.799637	48.794792
13.097208	20.28583	24.793974	14.110478	21.701134	22.217012
33.003544	21.11041	25.00658	19.122992	32.398567	13.605098
15.1145315	23.088867	27.474783	19.364998	26.487135	27.499458
28.697094	21.21718	18.703201	26.775208	14.010719	21.692347
18.372562	43.11582	29.081839	20.289959	23.680176	18.308306
17.204844	18.320065	24.393475	26.396057	19.094141	13.3019905
22.15311	22.185797	8.516214	18.894428	21.792608	19.331121
18.197924	7.5006843	22.406403	20.004215	14.412416	22.503702
28.53306	21.591028	13.810223	20.497831	21.898977	23.104464
49.99585	16.242056	30.294561	50.001595	17.771557	19.053703
10.399217	20.378187	16.49973	17.183376	16.70228	19.495337
30.507633	28.98067	19.528809	23.148346	24.391027	9.521643
23.886024	49.995125	21.167099	22.597813	19.965279	13.4072275
19.948694	17.087479	12.738807	23.00453	15.222122	20.604322
26.207253	18.09243	24.090246	14.105	21.689667	20.08065
25.010437	27.874954	22.92366	18.509727	22.190847	24.004797
14.788686	19.89675	24.39812	17.796036	24.556297	31.970308
17.774675	23.356768	16.134794	13.009915	10.98219	24.28906
15.56895	35.209793	19.605724	42.301712	8.797891	24.400295
14.086652	15.408639	17.301126	22.127419	23.09363	44.79579
17.776684	31.50014	22.835577	16.888603	23.925127	12.097476
38.685944	21.388391	15.98878	23.912495	11.909485	24.960499
7.2018585	24.696215	18.201897	22.489008	23.03332	24.260433
17.101519	17.805563	13.493165	27.105328	13.311978	21.913465
20.00738	15.405392	16.595737	22.301016	24.708412	21.422579
22.878702	29.606575	21.877811	19.900253	29.605219	23.407152
13.781474	24.454706	11.897682	7.2203646	20.521074	9.725295
19.30007	25.10501	11.688618	17.1001722	14.1800284	28.618876

#R Squared error

```
score = metrics.r2_score(y_train,training_data_prediction)
```

```
print(score) #This value should close to 1
```

```
0.9999948236320982
```

#Mean Absolute error

```
score2=metrics.mean_absolute_error(y_train,training_data_prediction)
```

```
print(score2)
```

```
0.0145848437110976
```

```
#prediction on test data
```

```
test_data_prediction=model.predict(x_test)
```

```
score3 = metrics.r2_score(y_test,test_data_prediction)
```

```
score4= metrics.mean_absolute_error(y_test,test_data_prediction)
```

```
score3
```

```
0.8711660369151691
```

```
score4
```

```
2.2834744154238233
```


```
#Visualization
```

```
plt.scatter(y_train,training_data_prediction)
```

```
plt.xlabel("Actual Price")
```

```
plt.ylabel("Predicted Price")
```

```
plt.show()
```



✓ 0s completed at 12:54

● ×